

Supplementary Information

Ontology-constrained multi-LLM scoring of hypothesis support in the predictive processing literature


Hamed Nejat, Jacob A. Westerberg, Alexander Maier, Jesse Spencer-Smith, and André M. Bastos

Supplementary Information

Supplementary Figures

Supplementary Fig. S1


Supplementary Fig. S1: Evaluation Pipeline Architecture. The evaluation pipeline flows from left to right in two independent streams (A to B to C and 1 to 2 to 3). The A to B to C pathway establishes the instructions, canonical glossary, and prompt for the models. The 1 to 2 to 3 pathway ingests the paper and combines figures with text to provide a single text object for the models. Steps 5 and 6 loop between models and papers until all evaluations are completed. The final step 7 quantifies and visualizes the evaluations.



Supplementary Fig. S2

Supplementary Fig. S2: Pairwise 2D Hypothesis-Space Projections and Human–Council Concordance. Panels show paper-level hypothesis scores for Local Oddball (LO; left column) and Global Oddball (GO; right column): (A1) LO $H_1 \times H_2$ ($N = 20$); (B1) GO $H_1 \times H_2$ ($N = 16$); (A2) LO $H_1 \times H_3$ ($N = 10$); (B2) GO $H_1 \times H_3$ ($N = 7$); (A3) LO $H_2 \times H_3$ ($N = 12$); and (B3) GO $H_2 \times H_3$ ($N = 10$). Panel N denotes the intersection of papers evaluated by both human experts for both hypotheses shown; per-hypothesis sample sizes are reported in Supplementary Table S14. Blue circles denote Council paper scores and red triangles denote two-expert human-consensus scores. Large gold-bordered markers show panel-specific means $\pm$ SEM across the same independent papers. Dotted gray lines connect paired evaluations for the same study. For comparison, the single-expert evaluation of Westerberg & Xiong (2025; open marker 30h*) is plotted against its corresponding Council evaluation, demonstrating the selective low-GO outlier displacement.


Supplementary Fig. S2: Pairwise Hypothesis Concordance across Independent Literature Studies (3x2 Projection)
Autonomous Council (Blue) vs. Human Domain Consensus (Red) with Panel-Specific Grand Means ($\pm$ SEM)



Supplementary Fig. S3


Supplementary Fig. S3: Decoding Robustness and Stochastic Replicate Reliability Analysis. (A) Cross-temperature score stability in the primary complete evaluation model (gemma-4-31b-it; 279 calls across 31 papers), showing Pearson correlations ($r \geq 0.9760$) and mean absolute errors ($MAE \leq 0.0421$) across sampling temperatures $T \in \{0.00, 0.35, 0.70\}$.

(B) Replicate reliability across independent stochastic generations held at fixed temperatures ($T = 0.00 : \bar{r} = 0.9905, \text{MAE} = 0.0126$; $T = 0.35 : \bar{r} = 0.9724, \text{MAE} = 0.0384$; $T = 0.70 : \bar{r} = 0.9444, \text{MAE} = 0.0585$), confirming that scoring variability across independent stochastic runs remains tightly bounded relative to cross-paper literature variance.



Supplementary Fig. S4


Supplementary Fig. S4: Evidence Grounding Control and Shuffled Text Sensitivity Analysis. Evaluates 54 inference calls across intact manuscript text (FULL), shuffled-evidence text (H_0), and prior-only prompts without paper text (PRIOR). Intact text yields strong differentiation between local and global contexts, whereas shuffled evidence and prior-only inputs collapse toward parity (0.450–0.493), showing sensitivity to supplied empirical evidence beyond prior-only input.



Supplementary Fig. S5



Supplementary Fig. S5: Figure-Text-Removal Sensitivity Analysis (Full 31-Paper Corpus Ablation). (A) Paper-level mean score correlation across the corpus between intact

inputs (FULL) and figure-text-removed inputs (NO_FIGURE; Spearman $\rho = 0.4718$ across $N = 28$ paired empirical paper means; factor-cell level MAE = 0.2253 across all 533 paired cells nested in 31 papers). (B) Factor-cell level directional concordance and stability across the corpus (directional concordance = 99.42% on non-zero cells [518/521], all-cell sign agreement = 97.19% [518/533], Pearson $r = 0.8539$), demonstrating that while removing figure descriptions produces measurable paper-level score adjustments, overall factor directionality was preserved across the literature corpus, while paper-level rank correspondence was moderate.



Supplementary Fig. S6

Supplementary Fig. S6: Evaluator Concordance Strata and Single-Study Sensitivity Analysis. (A) Mean squared difference (MSD) across evaluator strata ($K = 2$ human domain experts and autonomous LLM council) compared against a null permutation reference distribution (H_0 , $B = 10,000$ iterations). Agreement within human experts ($H-H = 0.1779$) and between human consensus and council ($H-A = 0.1431$) is significantly tighter than chance ($H_0 = 0.7490$, $p_{\text{perm}} = 0.0001$, 0/10,000 extreme permutations). (B) Council condition means for Local Oddball (LO) and Global Oddball (GO) across the entire corpus ($N = 31$), excluding Westerberg & Xiong (2025; $N = 30$), and for Westerberg & Xiong alone ($N = 1$), demonstrating that the aggregate corpus contrast is robust across published studies and that the preprint represents a selective low-GO outlier.



Supplementary Tables

#	Factor Name	Definition	Scope
1	Subtractive Inhibition (SST)	Linear input suppression via somatic inhibition (Somatostatin-mediated).	LO+GO
2	Divisive Inhibition (PV)	Multiplicative gain reduction (Parvalbumin-mediated).	LO+GO
3	Inhibition (GABA)	Chloride-mediated inhibition for prediction suppression.	LO+GO
4	Habituation to Sequence	Habituation suppresses local oddball response relative to a novel local oddball.	LO
5	Synaptic Depression (Adaptation)	Passive fatigue of synaptic efficacy due to repetition of stimulus.	LO
6	Activity Suppression	Reduction in firing rates for expected (predictable) stimuli.	LO+GO
7	Selective Sharpening	Signal-to-noise increase for predictable stimuli, where noise is selectively suppressed to highlight relevant information.	LO+GO
8	Alpha/Beta Mediated Suppression	Low-frequency bands (8–30Hz) associated with predictions, inhibiting prediction error signals.	LO+GO
9	VIP-Mediated Disinhibition	VIP-mediated inhibition of SST/PV neurons, leading to disinhibited pyramidal activity during prediction error.	LO+GO
10	Precision Weighting (Gain)	Top-down amplification of error units, leading to disinhibited gain in response to attended or highly precise stimuli.	LO+GO
11	E/I Balance Shift	Dynamic adjustment toward higher inhibition for predictable stimuli, leading to suppressed neural firing.	LO+GO
12	Omission Response	A generated internal signal due to absence of expected input.	LO+GO

Supplementary Table S1: The factors of Hypothesis 1 (Predictive Suppression). LO = local oddball. GO = global oddball.

#	Factor Name	Definition	Scope
13	Error Unit Depolarization	Supragranular pyramidal excitation driving ascending error signals.	LO+GO
14	Gamma Oscillations (30–90Hz)	High-frequency rhythms signaling ascending prediction errors.	LO+GO
15	NMDA Receptor Activation	Nonlinear amplification of error signals in superficial cortical layers.	LO+GO
16	Feedforward Flow	Supragranular-to-granular/supragranular ascending anatomical projection of errors.	LO+GO
17	Top-Down Error Gating	Feedback modulation of ascending prediction error channels.	LO+GO
18	MMN-like Latencies	Early sensory mismatch latencies (100–200ms) characteristic of local error computation.	LO
19	P300-like Latencies	Late cognitive mismatch latencies (300–500ms) characteristic of global contextual error computation.	GO
20	Inter-Laminar Error Divergence	Dissociation of error processing between granular and infragranular layers.	LO+GO
21	Phase-Amplitude Coupling	Theta/alpha phase modulation of gamma error bursts during surprise.	LO+GO
22	Thalamic Error Routing	First-order vs higher-order thalamocortical routing of sensory prediction mismatches.	LO+GO
23	Predictive Error Attenuation	Dampening of ascending prediction errors following successful model updating.	LO+GO
24	State-Dependent Error Routing	Behavioral state and arousal modulation of feedforward error gain.	LO+GO

Supplementary Table S2: The factors of Hypothesis 2 (Feedforward Prediction Errors).

#	Factor Name	Definition	Scope
25	Cross-Modal Ubiquity	Implementation of predictive coding mechanisms across multiple sensory modalities (visual, auditory, somatosensory).	LO+GO
26	Laminar Invariance of Principles	Preservation of predictive coding laminar logic across sensory, motor, and association cortices.	LO+GO
27	Cross-Species Homology	Conservation of predictive mechanisms across rodent, non-human primate, and human neocortex.	LO+GO
28	Scale-Free Predictive Dynamics	Fractals and hierarchical self-similarity of predictive computations across temporal scales.	LO+GO
29	Universal Error Coding	Common computational principles for error generation irrespective of receptive field complexity.	LO+GO
30	Canonical Microcircuit Generalization	Universal deployment of the canonical cortical microcircuit for inference across areas.	LO+GO
31	Cortico-Subcortical Homology	Shared predictive coding architecture between neocortex and basal ganglia / cerebellum.	LO+GO
32	Neuromodulatory Invariance	Consistent cholinergic/noradrenergic precision tuning across hierarchical stages.	LO+GO
33	Temporal Stability of Ubiquity	The consistent and non-transient presence of these mechanisms across the hierarchy over long recording sessions.	LO+GO
34	Hierarchical Order Stability	Evidence that the predictive mechanism is ubiquitous and not localized to a single hierarchical pole.	LO+GO
35	Population-Wide Ubiquity	Consistency of the mechanism across diverse neural populations and cell types throughout the hierarchy.	LO+GO
36	State-Independent Ubiquity	Presence of the mechanism across different brain states (e.g., wakefulness, sleep, anesthesia) throughout the hierarchy.	LO+GO

Supplementary Table S3: The factors of Hypothesis 3 (Ubiquity).

Term / Context	Description	Definition (notation)
Oddball	A surprising event, often a mismatch	$E = \{B\}$ given $S = \{A\}$
Local Oddball (LO)	An oddball that is immediate or short-term, often confounded by low-level adaptation mechanisms.	$E = \{AB\}$ given $S = \{AA\}$
Global Oddball (GO)	An oddball that is not LO, often contextual, that is not caused by adaptation mechanisms.	$E = \{AA\}$ given $S = \{AB\}$

Supplementary Table S4: Context definitions and formal expectation notations.

#	Study Identifier	Modality / Preparation	Inclusion Basis	Bibliographic Citation
1	Attinger et al., 2017	Mouse V1 (2-Photon Ca^{2+})	Comparative Set (Visuomotor Mismatch)	<i>Cell</i> 169, 1291–1302
2	Bakhtiari et al., 2021	In Silico (Deep Neural Nets)	Comparative Set (Hierarchical Streams)	<i>Nat. Commun.</i> 12, 25164
3	Bastos et al., 2012	In Silico / Theory (Microcircuit)	Comparative Set (Canonical Microcircuit)	<i>Neuron</i> 74, 695–708
4	Bastos et al., 2020	Primate V1/V4/PFC (Laminar ECoG)	Comparative Set (Laminar Rhythms)	<i>Neuron</i> 108, 124–137
5	Bekinschtein et al., 2009	Human (Scalp EEG)	Comparative Set (LO vs GO Originator)	<i>PNAS</i> 106, 1672–1677
6	Chao et al., 2018	Primate Temporal/Frontal (ECoG)	Comparative Set (Large-Scale Hierarchy)	<i>Nat. Commun.</i> 9, 3188
7	Friston et al., 2010	Theory (Free Energy Principle)	Canonical Extension (Bayesian Inference)	<i>Nat. Rev. Neurosci.</i> 11, 127
8	Furutachi et al., 2024	Mouse V1 (2-Photon / Opto)	Comparative Set (VIP/SST Disinhibition)	<i>Nature</i> 633, 398–406
9	Garrett et al., 2020	Mouse V1 (2-Photon Ca^{2+})	Comparative Set (Novelty vs Deviance)	<i>eLife</i> 9, e50340
10	Greedy et al., 2022	In Silico (Predictive Coding SNN)	Comparative Set (Error Back-propagation)	<i>NeurIPS</i> 35, 24213–24225
11	Hertag et al., 2020	In Silico (Spiking Interneurons)	Comparative Set (SST/PV Dynamics)	<i>eLife</i> 9, e57541
12	Jiang et al., 2024	In Silico (Energy-Based Model)	Comparative Set (Sequence Prediction)	<i>PLOS Comput. Biol.</i> 20, e1011801
13	Keller et al., 2012	Mouse V1 (2-Photon Ca^{2+})	Comparative Set (Sensorimotor Mismatch)	<i>Neuron</i> 74, 809–815
14	Keller et al., 2018	Synthesis (Cortical Circuits)	Comparative Set (Circuit Review)	<i>Neuron</i> 100, 424–435
15	Kiebel et al., 2008	In Silico (Hierarchical Dynamics)	Comparative Set (Temporal Hierarchy)	<i>PLoS Comput. Biol.</i> 4, e1000209
16	Lao et al., 2023	Human (Auditory Scalp EEG)	Comparative Set (Omission Response)	<i>Sci. Adv.</i> 9, eabq8657
17	Lee et al., 2025	In Silico (Multi-Area Network)	Comparative Set (Primate Oscillations)	<i>PLOS Comput. Biol.</i> 21, e1013469
18	Mikulasch et al., 2023	In Silico / Theory (Plasticity)	Comparative Set (Dendritic Error Review)	<i>Trends Neurosci.</i> 46, 45–59
19	Nejad et al., 2025	Primate / In Silico (Laminar Model)	Comparative Set (Layered Routing)	<i>J. Neurosci.</i> 45, 1234
20	Payeur et al., 2021	In Silico (Dendritic Plasticity)	Comparative Set (Burst Multiplexing)	<i>Nat. Neurosci.</i> 24, 1010–1019
21	Rao et al., 2024	Synthesis (Sensorimotor Neocortex)	Canonical Extension (25-Year Synthesis)	<i>Nat. Neurosci.</i> 27, 1221–1235
22	Rao & Ballard, 1999	In Silico (Hierarchical Model)	Canonical Extension (Foundational Model)	<i>Nat. Neurosci.</i> 2, 79–87
23	Sacramento et al., 2018	In Silico (Dendritic Network)	Comparative Set (Somatodendritic Model)	<i>NeurIPS</i> 31, 8721–8731
24	Spratling et al., 2008	In Silico (Biased Competition)	Comparative Set (PC-BC Synthesis)	<i>Biol. Cybern.</i> 98, 283–292
25	Spratling et al., 2010	In Silico (V1 Microcircuit)	Comparative Set (V1 Microcircuit Model)	<i>J. Neurosci.</i> 30, 3531–3543
26	Srinivasan et al., 1982	Fly Retina (Intracellular)	Canonical Extension (Redundancy Baseline)	<i>Proc. R. Soc. Lond. B</i> 216, 427
27	VanDerveer et al., 2021	Human (Clinical EEG Cohort)	Comparative Set (Sensory Gating Deficits)	<i>Schizophr. Bull.</i> 47, 1385
28	Wacongne et al., 2011	Human (Auditory MEG)	Comparative Set (Hierarchical Rule MMN)	<i>PLoS Biol.</i> 9, e1001189
29	Wacongne et al., 2012	Human EEG & Spiking Network	Comparative Set (Auditory Oddball Model)	<i>J. Neurosci.</i> 32, 3665–3678
30	Westerberg & Xiong, 2025	Primate V1/V4/IT (Laminar MUA)	Primary Index Study (MaDeLaNe)	<i>bioRxiv</i> 10.1101/2024.10.02
31	Yamins et al., 2014	In Silico / Primate IT (HCNN)	Comparative Set (Feedforward Baseline)	<i>PNAS</i> 111, 8619–8624

Supplementary Table S5: 31-Paper Literature Corpus Registry with modality, inclusion basis, and provenance.

Model Name	Parameters	Architecture & Provenance
gemma-3-27b-it	27B	Google DeepMind (Dense Transformer)
gemma-4-31b-it	31B	Google DeepMind (Extended Reasoning Architecture)
deepseek-r1-distill-llama-70b	70B	DeepSeek-AI / Meta Llama Distillation
gpt-oss-claude-4.5-sonnet	20B	Community Distillation / Reasoning
mistral-nemo-12b-thinking	12B	Mistral AI & NVIDIA
olmo-3-32b-think	32B	Allen Institute for AI (Ai2)
phi-4-reasoning-plus	14B	Microsoft Research
qwen3-14b-gemini-3-pro	14B	Alibaba Cloud / Community Fine-tune
qwen3-14b-claude-4.5-sonnet	14B	Alibaba Cloud / Community Fine-tune
qwen3.5-40b-claude-4.5-opus	40B	Alibaba Cloud / High-Capacity MoE

Supplementary Table S6: Open-weight LLM evaluator models evaluated in this work.

Evaluator Comparison	Level	Paired N	Pearson r	p	MAE	Concordance
Human–Human (H_1 vs H_2)	Factor cell	299	0.390	< 0.0001	0.308	88.96%
Human Consensus vs LLM Council	Study level	25	0.529	0.0125	—	—
Human Consensus vs LLM Council	Factor cell	175	0.162	0.0322	0.301	93.14%
Council Inter-Model Average	Factor cell	584	0.540	< 0.0001	0.284	85.00%
Shuffled Evidence (Null H_0)	Factor cell	191	-0.087	0.2314	0.662	56.54%

Supplementary Table S7: Evaluator Strata Agreement and Error Hierarchy Summary ($K = 2$ independent human domain experts). The independent study-level comparison ($N = 25$ papers with jointly evaluated factor ratings; Pearson $r = 0.529$, $p = 0.0125$, 95% CI [0.169, 0.764]) treats the independent paper as the unit of analysis. Factor-cell comparisons are descriptive summaries across scored cells nested within papers. Directional concordance denotes agreement in sign on non-zero factor cells.

Adjudication Category	Cells (N)	Proportion	Primary Disagreement / Failure Mode
Directly supported	6	20.0% (6/30)	Full primary text, equation, or figure corroboration of Council assignment
Partially supported	18	60.0% (18/30)	Relevant primary evidence present; conservative magnitude attenuation or specificity nuance
Unsupported attribution	6	20.0% (6/30)	Subtype gain conflation ($N = 3$), laminar error inversion ($N = 1$), untested modality penalty ($N = 1$), cross-paper attribution ($N = 1$)
Total Audited Disagreements	30	100.0% (30/30)	Diagnostic audit of $H - C \geq 0.50$ cells ($N = 296$ eligible paired cells)

Supplementary Table S8: Diagnostic Primary-Source Audit across 30 Numerically Selected Large Human–Council Disagreement Cells ($|H - C| \geq 0.50$ among 296 eligible paired cells; 10.1% of paired universe). Evaluates whether large disagreements reflect unsupported evidence attribution or differences in magnitude, contextual interpretation, or biological specificity despite relevant primary-source evidence. Because this diagnostic sample was enriched for large discrepancies, these proportions characterize disagreement modes and do not represent corpus-wide accuracy or hallucination rates.

Decoding Setting / Comparison	Evaluated Units (N)	Correlation	MAE	Paper Ranking Stability
Temperature (gemma-4-31b-it)	Stability			
$T = 0.00$ vs. $T = 0.35$	679 factor cells	Pearson $r = 0.9900$	0.0266	Spearman $\rho_{\text{rank}} = 0.9596$
$T = 0.00$ vs. $T = 0.70$	679 factor cells	Pearson $r = 0.9760$	0.0421	Spearman $\rho_{\text{rank}} = 0.9212$
$T = 0.35$ vs. $T = 0.70$	679 factor cells	Pearson $r = 0.9816$	0.0367	Spearman $\rho_{\text{rank}} = 0.9482$
Replicate (gemma-4-31b-it)	Reliability			
Replicates at $T = 0.00$	639 triples	Mean $\bar{r} = 0.9905$	0.0126	Spearman $\rho_{\text{rank}} \geq 0.9818$
Replicates at $T = 0.35$	619 triples	Mean $\bar{r} = 0.9724$	0.0384	Spearman $\rho_{\text{rank}} \geq 0.9782$
Replicates at $T = 0.70$	572 triples	Mean $\bar{r} = 0.9444$	0.0585	Spearman $\rho_{\text{rank}} \geq 0.9745$

Supplementary Table S9: Factorial Robustness Analysis across Sampling Temperatures and Stochastic Replicates for the Complete Evaluation Model (**gemma-4-31b-it**, 279/279 factorial conditions). Quantitative factorial statistics are reported exclusively for **gemma-4-31b-it** because it completed full factorial coverage across all 31 papers $\times$ 3 temperatures $\times$ 3 repeats. Extended chain-of-thought models (**phi-4-reasoning-plus** and **olmo-3-32b-think**) completed 130/279 and 30/279 calls, respectively, but frequently exhausted dynamic generation token budgets during chain-of-thought reasoning prior to emitting the terminal JSON scoring block, precluding equivalent comparative factorial statistics. Paper ranking stability measures Spearman correlation between paper mean scores across conditions ($N = 29$ evaluable papers).

Evidence Manipulation	Calls (N)	Mean LO Score [$\pm$ SD]	Mean GO Score [$\pm$ SD]
Intact Evidence (FULL)	18	0.512 ± 0.081	0.474 ± 0.076
Shuffled Evidence (H_0)	18	0.489 ± 0.092	0.493 ± 0.088
Prior Only (Zero Evidence)	18	0.450 ± 0.110	0.450 ± 0.108

Supplementary Table S10: Evidence Grounding vs. Prior Bias Sensitivity Analysis (54 Calls across Intact, Shuffled, and Prior-Only Controls). Under intact text, models distinguish Local vs. Global Oddball contexts, whereas under shuffled text (H_0) and prior-only prompts without manuscript text, scores across contexts collapse to parity.

Ablation Condition	Paired Cells (N)	Pearson r	MAE	Directional Concordance
FULL Baseline (With Figures)	533	1.000	0.000	100.0% (Baseline)
NO_FIGURE (Figures Stripped)	533	0.8539	0.2253	99.42% (518/521 non-zero cells)

Supplementary Table S11: Figure-Text-Removal Sensitivity Summary ($N = 31$ papers).

Dimension	Stratum	Papers (N)	Council LO Mean	Council GO Mean
Study Type	Empirical	13	0.419	0.428
Study Type	Computational / Theory	18	0.351	0.365
Modality	In Silico	19	0.363	0.376
Modality	Human In Vivo	7	0.357	0.416
Modality	Animal In Vivo	5	0.511	0.424
Publication Era	1982–2014	12	0.368	0.393
Publication Era	2015–2020	7	0.447	0.488
Publication Era	2021–2025	12	0.347	0.309
Publication Status	Peer-Reviewed Journal	30	0.381	0.392
Publication Status	Preprint (Westerberg & Xiong 2025)	1	0.453	0.053

Supplementary Table S12: Corpus Subgroup Sensitivity Analysis across Taxonomic Strata. Demonstrates that council-estimated evidence support across Local and Global Odd-ball contexts varies by experimental modality and publication era.

Validation Axis	Evaluated (N)	Strictly Correct	Partial / Minor	Audit Pass Rate
Axis A: Numerical Trend & Data Extraction	10	9	1 (Minor phrasing)	100.0% PASS (10/10)
Axis B: Empirical Effect Direction	10	10	0	100.0% PASS (10/10)
Axis C: HPC-36 Ontology Factor Mapping	10	9	1 (Invertebrate)	100.0% PASS (10/10)
Overall Manual Validation Outcome	10 Figures	10/10	0 Fatal Inversions	100.0% PASS

Supplementary Table S13: Manual Expert Validation Summary of Generated Multimodal Figure Descriptions ($N = 10$ representative figures audited against primary source PDFs). All 10 figures met the prespecified audit criterion (Audit Pass Rate: 100.0%); minor phrasing and partial ontology mappings are reported separately.

Context	Hyp.	N	Human Mean	Council Mean	Human–Council [95% CI]	$t(df), p$
Local Oddball (LO)	H_1	21	$+0.537 \pm 0.041$	$+0.337 \pm 0.027$	$+0.200$ [0.111, 0.289]	$t(20) = 4.69, p = 0.0001$
	H_2	23	$+0.552 \pm 0.045$	$+0.363 \pm 0.026$	$+0.189$ [0.117, 0.261]	$t(22) = 5.42, p < 0.0001$
	H_3	14	$+0.520 \pm 0.034$	$+0.268 \pm 0.045$	$+0.252$ [0.161, 0.344]	$t(13) = 5.95, p < 0.0001$
Global Oddball (GO)	H_1	18	$+0.600 \pm 0.048$	$+0.299 \pm 0.031$	$+0.302$ [0.182, 0.422]	$t(17) = 5.32, p < 0.0001$
	H_2	22	$+0.560 \pm 0.043$	$+0.259 \pm 0.049$	$+0.300$ [0.205, 0.396]	$t(21) = 6.54, p < 0.0001$
	H_3	11	$+0.557 \pm 0.065$	$+0.312 \pm 0.057$	$+0.244$ [0.080, 0.408]	$t(10) = 3.32, p = 0.0077$

Supplementary Table S14: Primary Paper-Level Hypothesis Calibration and Concordance ($K = 2$ Human Consensus vs. Autonomous Council). Analysis is restricted to independent papers where both human experts ($K = 2$) evaluated mechanisms within the specific hypothesis. Council models were averaged within paper first. SEM denotes standard error across independent papers ($SD/\sqrt{N}$). Paired differences (Human–Council mean difference) evaluate systematic score attenuation. Paired t -tests ($t(df)$) and two-sided p -values evaluate the difference in mean support.

Omitted Evaluator Model	Paired N	Pearson r	p	Spearman ρ	p	MAE
deepseek-r1-distill-llama-70b	1710	0.9553	< 0.0001	0.9589	< 0.0001	0.0465
gemma-3-27b	1733	0.9747	< 0.0001	0.9472	< 0.0001	0.0408
gpt-oss-claude-4.5-sonnet	1801	0.9920	< 0.0001	0.9844	< 0.0001	0.0086
mistral-nemo-12b-thinking	1775	0.9732	< 0.0001	0.9698	< 0.0001	0.0239
olmo-3-32b-think	1798	0.9900	< 0.0001	0.9800	< 0.0001	0.0122
phi-4-reasoning-plus	1778	0.9825	< 0.0001	0.9735	< 0.0001	0.0255
qwen3-14b-gemini-3-pro	1800	0.9797	< 0.0001	0.9649	< 0.0001	0.0243
qwen3.5-40b-claude-4.5-opus	1724	0.9497	< 0.0001	0.9178	< 0.0001	0.0440

Evaluator profiles with hypothesis-level scores but unassigned individual factor cells:

gemma-4-31b	1816	—	—	—	—	—
gpt-oss-safeguard-120b	1816	—	—	—	—	—

Supplementary Table S15: Leave-One-Model-Out Council Score Stability Analysis across Evaluator Architectures. Evaluates the sensitivity of aggregate Council scores to single-model removal by comparing the complete Council continuous mean against the Council recomputed after omission of each contributing evaluator across matched factor cells ($N = 1710$ – 1801). Across all active evaluator omissions ($N = 8$ model omissions), Council scores remained highly stable: mean Pearson $r = 0.9747$ (95% CI: [0.9619, 0.9874], range 0.9497–0.9920), mean Spearman $\rho = 0.9621$ (95% CI: [0.9441, 0.9800], range 0.9178–0.9844), and mean MAE = 0.0282 (95% CI: [0.0163, 0.0402], range 0.0086–0.0465). Confidence intervals are computed using Student’s t -distribution ($df = N_{\text{omissions}} - 1 = 7$) with model

omission as the unit. MAE is reported in original $[-1.0, +1.0]$ score units. Two evaluator profiles contributed hypothesis-level Council scores but no individual factor-cell scores and therefore were not evaluable in this factor-level leave-one-out analysis.